

THE LANGUAGE OF SKETCHES

Xiaoyu Zhang

Carnegie Mellon University

Pittsburgh, PA 15213

April 2022

Thesis Committee

Oliver Kroemer, advisor, chair

Yonatan Bisk, advisor, co-chair

Jean Oh

*Submitted in partial fulfillment of the requirements for
the degree of Master of Science in Computer Science.*

Abstract

Creative AI has seen much progress in recent years. Works like DALL-E 2 can generate inspiring art pieces from text descriptions. Instead of synthesizing realistic art works from language, we approach creativity from a different angle and investigate composition of semantic parts and visual concepts in sketches. For example, people can draw a circle to represent the moon, a scoop of ice-cream, or the face of a cat. In a similar way, language descriptors can be composed to create new concepts. People can draw a large round cat face or a narrow oval cat face.

In order to study this reuse of abstract concepts, we construct a dataset of language annotated sketches. We examined current sketch datasets and found that they either lack language annotations or semantic part annotations. Therefore, we collect a dataset of 11,150 (sketch part, text) pairs for 572 face sketches and 787 angel sketches.

To understand the limits of current vision-language models, we fine-tuned CLIP, a model pre-trained with a contrastive objective on 400 million (image,text) pairs and can map (image,text) pairs into a joint vision-language embedding space. We observed that (1) CLIP cannot easily generalize to an unseen category on the task of pairing sketches with their descriptions even though similar shapes and descriptions have occurred in training; (2) through fine-tuning, average cosine distance has increased between a pair of descriptors used by annotators to differentiate two sketches. With insights gained about how language and sketches interact in the CLIP embedding space, our aim is to facilitate research into models that can generate sketches in a part-based manner satisfying descriptions given by users of the pictures they have on their minds.

Acknowledgements

I first and foremost thank my advisors, Professor Oliver Kroemer and Professor Yonatan Bisk, for advising me on this project. I am very grateful that I had the opportunity to explore the creative sketching domain with them during my Master's. Through working with them on this project, I have learned how humans use language creatively to describe sketches, and I am excited about building systems that can interact seamlessly with humans in the future. I want to thank Professor Jean Oh for being on my committee.

I also want to thank my lab mates in the IAM and COMEDY lab (the name *COMEDY* was coined in April 2022 ). Thank you guys for testing the numerous beta versions of my data collection interface and providing feedback without which the official deployment would not have been as successful.

I thank my parents for encouraging me to pursue research topics that I enjoy. Thank you guys for listening to all my crazy ideas.

Contents

Abstract	ii
Acknowledgements	iii
1 Introduction	1
2 Related Work	4
2.1 Sketch Dataset	4
2.2 Collaborative Drawing	6
3 Data Collection	9
3.1 Prompt-Guided Sketch Text Dataset	10
3.1.1 Overview	10
3.1.2 Interface Design	10
3.1.3 Results	14
3.2 Contrasting Sketch Text Dataset	15

3.2.1	Overview	15
3.2.2	Interface Design	16
3.3	Dataset Summary	18
4	Modeling	33
4.1	Task Definition	33
4.2	Method	34
4.2.1	Contrastive Objective	34
4.2.2	Text Encoder	36
4.2.3	Image Encoder	38
4.2.4	Fine-Tuning CLIP	39
5	Results & Analysis	41
5.1	Classification Experiment	41
5.2	Word Similarity Experiment	43
6	Future Work	46
A	Designing the Requirement Section for the Prompt-Guided Sketch Text Dataset	48

Chapter 1

Introduction

Creative AI, such as using deep learning models to generate paintings and music, has been a popular research domain. Since creative activities are hallmarks of human intelligence, numerous works have attempted to replicate the creative process on machines. For example, Pharmako-AI (Allado-McDowell & Okojie, 2020) is a book co-written by K Allado-McDowell and GPT-3 (Brown et al., 2020) through exchanges between the human and the language model. Works like DALL-E, GLIDE, and DALL-E 2 tackle the problem of synthesizing images from short language descriptions, and these generative models have produced many imaginative and inspiring art pieces (Ramesh et al., 2021; Nichol et al., 2021; Ramesh et al., 2022). These work are often motivated by the vision to create machines that can interact with people and augment human creativity. Our work is driven by a similar vision: how can we build systems that can be creative like humans so that AI agents and people can participate in creative activities together and inspire each other.

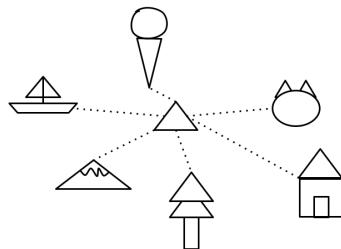


Figure 1.1: Triangle is adapted and composed with other shapes to create sketches of different objects. The objects are, from top and in clockwise order: *ice-cream*, *cat*, *house*, *pine tree*, *mountain*, and *boat*.

Instead of generating realistic art works from a wide variety of language like DALL-E, we approach creativity from a different angle and investigate composing basic visual concepts in object sketches. An example is illustrated in Figure 1.1: a triangle can be the sail of a boat, the cone of an ice-cream, the left/right ear of a cat, the roof of a house, the canopy of a tree, or the body of a mountain. The same triangle is adapted to become different parts in sketches of different objects. While sketches contain much fewer details compared to their natural images counterparts, people are not any less creative when sketching. One could say that the abstractness allows for more creativity.

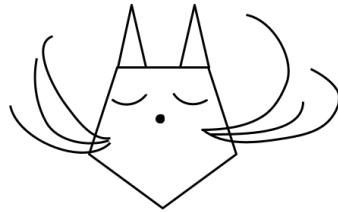


Figure 1.2: People are likely to provide a variety of descriptions when asked, “*What features does this cat sketch have?*”.

Similarly, language descriptors can be composed to form new concepts applicable to different objects. For example, combining *large* and *round* or *narrow* and *oval* to describe different kinds of cat faces. Many other objects can be *large round* or *narrow oval*: angel halo, glasses, mirror, plate, necklace charm, table, etc. Moreover, there are multiple ways to describe the same sketch, depending on the person’s viewpoint. *What visual features does this cat sketch have?* (Figure 1.2) Some people might pick up on its usual size; others might notice its slightly edged chin, and there are still others that focus on the long whiskers curving upwards. People are likely to provide a variety of responses.

We want to build systems that can compose basic shapes in a creative manner and understand how words are combined and adapted to describe different objects. In this thesis, we are interested in semantic parts in sketches and how people describe them. For example, a face sketch may contain four semantic parts: eyes, nose, mouth, and face contour, that can be drawn using various geometric shapes with different features. After examining existing sketch datasets, we found that they do not contain either language annotations or semantic part annotations. We explain the details of related sketch datasets in Chapter 2. To fill the gap and study this reuse of abstract concepts, we construct a dataset of language annotated face and angel sketches. We elaborate on the data collection process in Chapter 3.

Moreover, we hope to understand the limits of current vision-language models, such as CLIP, or Contrastive Language-Image Pre-Training, which is trained on millions of image-text pairs. CLIP has been used extensively to provide a loss function for optimizing text-to-image synthesis models

(Frans et al., 2021; Ramesh et al., 2022; Gal et al., 2021; Patashnik et al., 2021). Despite the abundance of knowledge captured by CLIP’s embedding space, we observed that CLIP, trained to associate natural images with their captions, has difficulty matching sketches from unseen categories with their part descriptions, suggesting that they might not understand that sketches from different categories are composed of parts different in semantics but share similar shapes and descriptions. We explain the details of CLIP and the fine-tuning process in Chapter 4, and we analyze CLIP’s performance and changes in its text embeddings after fine-tuning on our dataset in Chapter 5.

Chapter 2

Related Work

We survey two areas of related work: text-to-image synthesis and sketch representation learning. There is a long line of works on generating realistic images by learning their underlying distributions, especially on employing generative adversarial networks (GAN). Related to our work is a set of works on generating images from text descriptions or use texts to manipulate image styles, but they often work with a constraint set of language, with the exception of industrial-scale generative models with billions of parameters trained from text-image dataset on the hundred million scale. Most of these work focus on datasets of photo-realistic images, and our work instead centers around sketches, which are abstract since certain features of the objects in the sketches are illustrated figuratively and contain sparse information, as most part of the image is empty. Therefore, we also look into the field of sketch representation learning: what methods have been employed and what datasets are constructed.

2.1 Sketch Dataset

Since our work seeks to collect a dataset of (sketch parts, text description), we survey existing sketch datasets and find that they either lack language descriptions, for the entire sketch or for parts in a sketch, or they do not contain semantic object annotations.

TU-Berlin Eitz et al. (2012) is one of the first works to investigate the characteristics of free-hand sketches and attempt to extract local features based on orientation, which are later used in the task of sketch recognition. It also provides the TU-Berlin sketch dataset that contains 20,000 sketches spanning 250 object categories with 80 samples in each. The TU-Berlin dataset is then extended for sketch-based 3D object retrieval to contain 1814 new sketches for 130 common household object categories. This additional set also contains hierarchical category labels; for example, the *animal* category contains *arthropod*, *biped*, *human*, *flying creature*, *quadruped*, *underwater creature* categories if going one level down the hierarchy. The TU-Berlin contains some of the highest quality sketches for a wide range of categories, but they lack both text descriptions and semantic part annotations, and our dataset contains both types of annotations, although only for simpler sketches in the face and angel category.

Datasets with Text Descriptions The QMUL dataset contains sketches of shoes and chairs: 419 shoe sketches of various types of shoes, such as high heels, boots, and ballerina flats, and 297 chair sketches of different kinds (Yu et al., 2016). Although QMUL annotates for attributes of the sketches, they come from a fixed set of descriptions and are obtained from product tags, so these descriptions do not reflect how humans would creatively describe the drawings on their mind, which is a feature of our collected dataset. The Sketchy dataset contains 75,000 sketches of 125 categories, and each sketch is paired with an image; Sketchy is additionally annotated with sketch validity metrics and short descriptions from annotators (Sangkloy et al., 2016). However, these descriptions are more like comments given by the participants, and the collecting process is not carefully designed, unlike our data collection process, where obtaining the text description is the main focus, ensuring the quality of the language. Moreover, the short texts in Sketchy do not describe individual semantic object in the sketches.

Quick,Draw! Dataset Quick,Draw! collects the largest sketch datasets containing 50 million sketches distributed across 345 object categories, each containing around 100,000 sketch (Ha & Eck, 2017). Annotators for the TU-Berlin dataset have 30 minutes to draw one sketch, and annotators for the QMUL dataset have reference photos for the sketches they would draw later; on the other hand, Quick,Draw! participants had 20 seconds to create the sketches for a randomly assigned category without time to look for reference, so the Quick,Draw! sketches are the simplest; however, since sketches that failed to be recognized by a classification model are filtered, they are in general of good quality and are representative of the general population drawing skills. Moreover, many works on sketch representation learning use the Quick,Draw! dataset, so utilize these sketches allow us to potentially adapt some of the pre-trained generative models. These features of the Quick,Draw! dataset make it very favorable to be used to learn an collaborative drawing robot that can interact

with a wide range of users, including children. However, Quick,Draw! dataset contains neither semantic part annotation nor language descriptions, and our collected dataset contains both in order to build collaborative drawing robots.

Datasets with Semantic Part Annotations The sketches in our dataset come from the Quick,Draw! dataset, but the original Quick,Draw! dataset does not provide labels for semantic objects in the sketches. The Sketch Perceptual Grouping (SPG) dataset (Li et al., 2018) and the SketchSeg dataset (Qi & Tan, 2019) contain annotations for semantic segmentation, meaning that each stroke in a sketch is paired with a semantic label; for example, strokes in a sketch for the *face* category are assigned one of the following labels: *eyes, nose, mouth, ear, hair, moustache, outline of face*. The SPG dataset annotates for 25,000 Quick,Draw! sketches, belonging to 25 (out of 345) categories, and it selects 800 out of the original 100,000 sketches in each category for annotation. The SketchSeg dataset builds upon the SPG dataset by using a recurrent neural work to generate additional sketches stemmed from one sketch in SPG; since each stroke in the generated sketch corresponds to a stroke in the original sketch, it automatically has part annotations. However, since the quality of the generated sketch is dependent upon the generative model, we choose the original SPG dataset, which is also publicly available. The SPG dataset does not contain text descriptions like our dataset, so although they can be used to build a collaborative drawing system, such system would not have the capacity to respond to natural language commands from users, and our work is interested in learning collaborative robots that can communicate freely with people.

2.2 Collaborative Drawing

Sketch-RNN Along with the Quick,Draw! dataset, Ha & Eck (2017) introduces Sketch-RNN and provides a web demo of collaborative drawing, where users can select a category and draw the first few strokes of a sketch, and the model will complete the rest of the sketch¹. Sketch-RNN uses a variational autoencoder (VAE) with bi-directional RNN encoder and RNN decoder. Ha & Eck (2017) uses a vector format to represent the sketches: each sketch is a set of strokes, and each stroke is a sequence of points, so the smooth curve is approximated with piecewise linear splines created by connecting adjacent points in the sequence. Each point is a 5-component vector, $(\delta x, \delta y, p_1, p_2, p_3)$: the first two components represent changes in the (x, y) coordinate of the current point compared to the last one; the last three make up a one-hot vector representing current state of the sketch. Let $s = [p_1 \ p_2 \ p_3]$, $s = [1 \ 0 \ 0]$, if the current point will be connected with the next point;

¹You can access the demo from this website <https://magenta.tensorflow.org/sketch-rnn-demo>

$s = \begin{bmatrix} 0 & 1 & 0 \end{bmatrix}$, if the current point is the last point in the current stroke, so the pen is expected to be lifted next; $s = \begin{bmatrix} 0 & 0 & 1 \end{bmatrix}$, if the current point is the last point in the sketch. The RNN encoder takes this sequence as input and uses the final hidden state h of the RNN to parameterize a Gaussian distributions of size N_z , from which a latent random variable $z \in \mathbb{R}^{N_z}$ is sampled. The latent z will be used to (1) initialize the hidden state h_0 of the RNN decoder and (2) be concatenated with the vector at each time step to be fed as input into the decoder. The RNN decoder predicts parameters for distributions (Gaussian mixture model for sampling $\delta x, \delta y$ and categorical distribution for sampling p_1, p_2, p_3) at each time step from the hidden states of the RNN. Since the decoder works in an autoregressive manner, meaning that points in a sketch are passed one by one into the RNN, the decoder can *encode* the first few strokes provided by the users into the hidden vector, which is conditioned on for later steps to produce distributions of points that make up the rest of the sketch. Subsequent works on sketch representation learning can be grouped into those that represent the input in this stroke-point vector format and those that treat the entire sketch as a raster image.

Although our work uses the QuickDraw! sketches, we render the vectors into raster images in order to use CLIP and its vision-language joint embeddings, since compared to Sketch-RNN, we need to incorporate text descriptions of the sketch parts. It is challenging to align semantic information in text with distributions of the x, y coordinates. While language gives high-level guidance on how strokes are organized on canvas to illustrate a particular concept, the coordinate information is too low level if directly used without an intermediate level of abstraction to aggregate these low-level sequences.

DoodlerGAN DoodlerGAN (Ge et al., 2020) is a more recent work on collaborative drawing and represents sketches as raster images; it uses Generative Adversarial Network (GAN) to model the sketches². Compared to Sketch-RNN, which can only complete the missing parts of sketches once users are done with their parts, DoodlerGAN generates sketches using a part-based approach so that users and the system can take multiple turns to create sketches more collaboratively. Moreover, as indicated by Ge et al. (2020), compared to vector inputs, raster images can take better advantage of information from spatial structure of the parts and ignore effects of shifting coordinates by focusing more on relationships among the sketch parts. The part-based GAN has two networks: a part selector and a part generator. Given an incomplete sketch with some parts drawn (for example, a bird with its eyes and beak), the part selector determines what kind of part to draw next (for example, the bird wings). Then given the partial sketch and the category of the next part, the part generator produces a raster image of the part and its location on canvas. The part generator is a conditional GAN derived from StyleGAN2, and it takes as input an image with number-of-parts+1

²Try the DoodlerGAN demo here: <http://doodlergan.cloudcv.org/>

channels: one channel for each part and an additional channel for the whole partial sketch. The part selector uses the encoder of the part generator with an extra linear layer at the end as the classification head.

While the part-based generation aspect of DoodlerGAN aligns with our goal to create collaborative drawing agent, there is no language associated with each part, and the generation process is not guided by language supervision.

Chapter 3

Data Collection

To study the composition of visual concepts, we need to collect a sketch dataset that contain both semantic part and language annotations. We tested two ways of collecting such dataset. Initially, in order to create a dataset with creative composition of shapes, such as the use of triangle in different sketches illustrated in Figure 1.1, we designed text prompts like *happy angel* and asked annotators to draw sketches in response to the prompts while annotating each step in their drawings. After deployment of pilot, although we collected many creative sketches, we saw major problems related to misalignment between the drawn parts and the text descriptions. We explain this prompt-guided design and its associated problems in details of in Section 3.1.

After re-design, we leverage existing sketch datasets, QuickDraw and SPG (both explained in details in Section 2.1), to solve the problems of low-quality part and text annotations unusable for model learning. In the new version, we present two sketches at a time to annotators, where the sketches contain semantic parts that are visually contrasting to aid annotators in coming up with creative descriptions. We walk through the process of designing this *contrasting* design in Section 3.2. After receiving satisfactory pilot results, we deployed the design full scale on Amazon Mechanical Turk (AMT); we present statistics of the dataset and explain how it allows us to study creative composition of abstract concepts in Section 3.3.

3.1 Prompt-Guided Sketch Text Dataset

3.1.1 Overview

When we first started designing our data collection interface, we wanted to collect sketches for prompts similar to the ones used in DALL-E (Ramesh et al., 2021): creative composition of attributes and objects that are not commonly associated. For example, *an evil cup of bubble tea* and *happy moon*. In addition to sketches, we also require annotators to decompose their drawings into steps and provide descriptions for each step. We hope to discover combination of simple shapes like examples in Figure 1.1, and because the imaginative prompts would result in creative sketches, the part descriptions would also be interesting.

We deploy the data collection interface on Amazon Mechanical Turk (AMT), which is a crowd-sourcing website that hosts different machine learning annotation tasks. In the remaining text, we use the word *turker* to refer to annotators we recruit on AMT; we will also use the word *HIT*, Human Intelligence Task, to refer to a task hosted on AMT. Please refer to AMT FAQs for official definition and answers to questions related to AMT.

For each HIT, we need to design: (1) instruction and requirements explaining the dos and don’ts; (2) qualification task to train turkers to provide high-quality annotations; (3) the main interface for data collection. After deploying the pilot, we realized a few major problems with this design. Firstly, due to the subjective nature of sketching, although the sketches are very creative, it was hard to understand how some annotators are illustrating the given prompts. Moreover, turkers are taking more than 30 minutes for each task, and, most importantly, many descriptions do not align with the drawn objects, making the data difficult to use for model learning. For example, in one step, an annotator drew the entire cat face, but they only annotated *big eyes*.

3.1.2 Interface Design

Main Task Interface

We illustrate a typical annotation process in Figure 3.1. The annotator draws a step on the canvas, enter text description for this step in the *Annotation* column, and hit *Add* to display it as a new row in the annotation table. For the annotator’s convenience, we include an *Undo* button and a *Clear* button for erasing strokes and clearing the entire canvas. If the annotator wants to remove

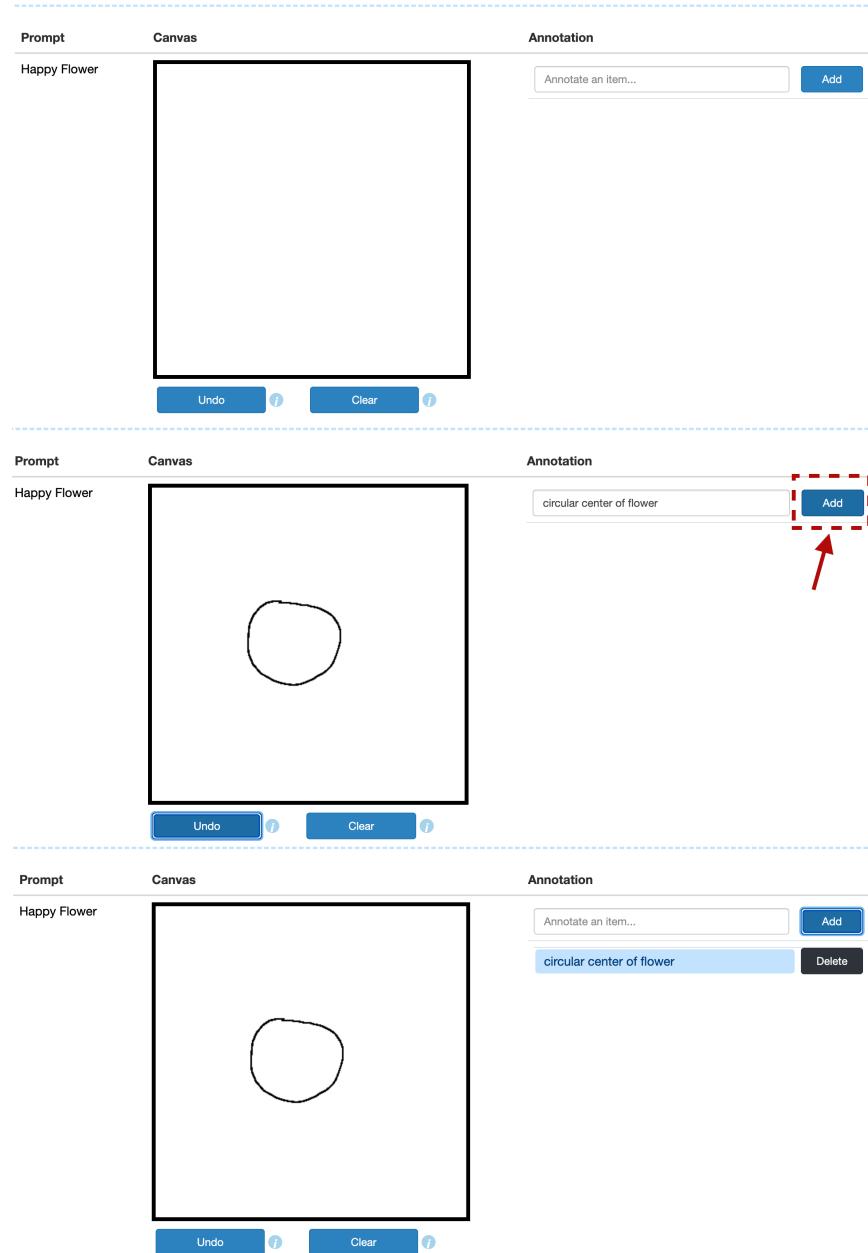


Figure 3.1: A typical annotation process. Top: interface at the start of annotation. Middle: before adding text descriptions for the drawing; red arrow and box show where to click to add text. Bottom: after adding text descriptions for the flower sketch.

an entire step, including the drawing and the text description, they can use the *Delete* button. An example is shown in Figure 3.2. Repeat the drawing-and-adding process until the drawing is done. This design encourages turkers to decompose their drawings into semantically meaningful parts.

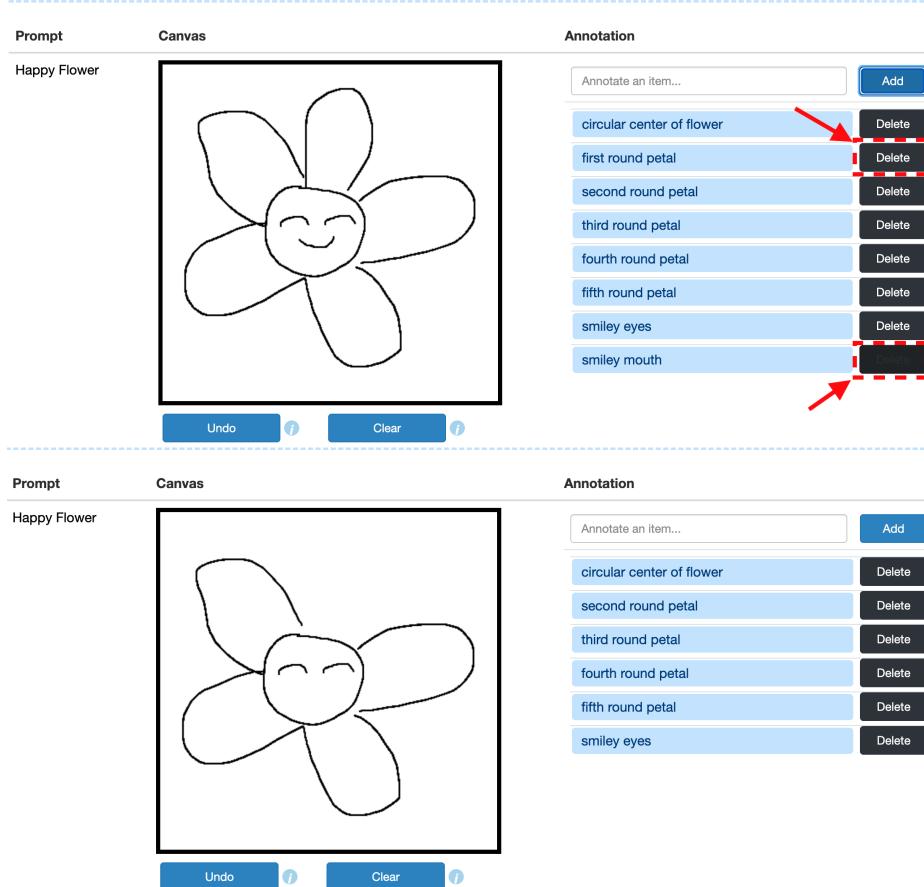


Figure 3.2: Top: the completed annotation for the prompt *Happy Flower*. Red arrows and boxes point to *Delete* buttons that can delete the text annotations along with their drawings. Bottom: after deleting the steps *first round petal* and *smiley mouth*.

We encountered some difficulties when implementing the *Delete* button. At the beginning, we treated erasing strokes as drawing the same strokes but in white; however, when strokes overlap each other, overwriting with white strokes breaks other strokes into segments. Therefore, we change the drawing canvas to use layers like Photoshop, so that deleting strokes would be the same as deleting an entire layer, leaving other strokes intact.

Instruction and Requirement

To ensure that turkers understand the purpose of collecting the dataset, the instruction begins with the motivation behind this project (Figure 3.3). What we struggled the most when drafting

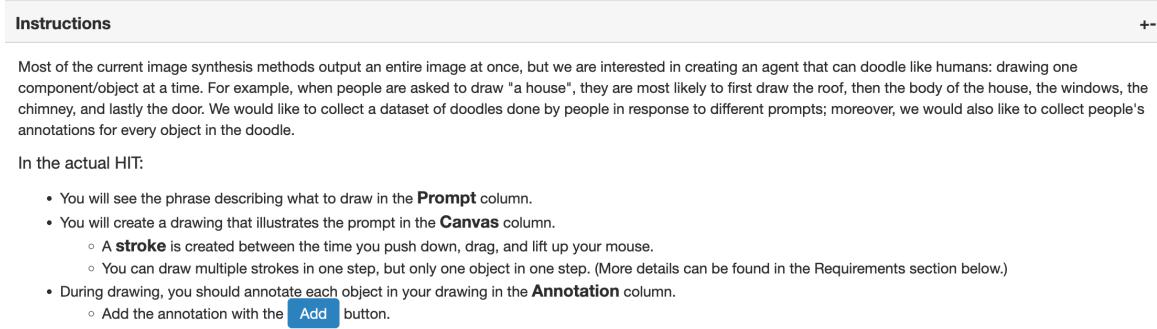


Figure 3.3: The instruction section used in the prompt-guided sketch text dataset.

the requirements was deciding what a single *step* in sketching was; how do we clearly explain this definition to the turkers? We considered providing a list of geometric shapes, such as rectangle, triangle, circle, etc., or asking annotations for each stroke, but these options could not reflect how people naturally sketch.

There is a wide spectrum of allowed annotations depending on how people sketch. For example, when drawing for the prompt *Happy Face*, one person might annotate 3 steps: *large u-shaped face, round eyes, big smiley mouth*. But for someone who likes to draw detailed eyes, they might describe the shape of the eye contour and the length of the eyelashes. The great variation in personal styles makes creative sketches fascinating to study but also challenging to collect a high-quality dataset.

We resorted to repeatedly testing the interface with lab mates to refine the requirements. The refinement process is explained in Appendix A. The requirements deployed in the final pilot is shown in Figure 3.4. Bad Example 2 is shown in Figure 3.5 as an instance of the examples used in the final requirements. To view all the examples, refer to: https://erinzhang1998.github.io/portfolio/amazon_anno.

Qualification

We set up a qualification test on AMT to (1) train turkers to have better understanding of the task and (2) to select turkers who can provide annotations that satisfy all the requirements. Similar to the process of writing the requirements, we went through several rounds of testing with students in the lab to come up with a set of questions correspond well with the requirements. The qualification test starts with the same instruction and requirements in the final HIT, thus allowing turkers to familiarize themselves with the requirements and ask clarification questions before completing the actual annotation task. The test leads with a navigation bar (Figure 3.6) to make it convenient for

Requirements

The entire drawing must satisfy:

1. **Each drawing should contain at least 2 steps.**
 - Do not complete the drawing in one step.
 - See details in: [Bad Example 1](#)

Every step must satisfy:

1. **Do not draw multiple objects in one step, unless they are multiple copies of the same object.**
 - For example, do not draw roof, window, and door of the house all in one step. *But*, you can draw two windows in one step.
 - See details in: [Bad Example 1](#) [Good Example 2 \(Step 3\)](#)
2. **Name the object in the annotation.**
 - For example, if you draw a roof, annotation should at minimum include the word "roof".
 - See details in: [Bad Example 2 \(Step 1\)](#) [Bad Example 2 \(Step 2\)](#)
3. **If you use the word "right", it should refer to this side: →. If you use the word "left", it should refer to this side: ←.**
 - For example, do not draw a right window of the house and then annotate it with "left window".
 - See details in: [Bad Example 2 \(Step 3\)](#) [Bad Example 2 \(Step 4\)](#)
4. **Differentiate between plural and singular.**
 - For example, if you draw two windows, annotation should include "windows" instead of "window".
 - See details in: [Bad Example 3 \(Step 3\)](#) [Good Example 2 \(Step 3\)](#)

Note:

1. **Not required to describe the geometric shape of the object, but preferred.**
 - For example, if you draw a triangle as the roof of a house, you do not need to annotate "triangular roof"; annotating with "roof" is okay, since this satisfies "Name the object in the annotation".
 - We do not require describing the geometric shape, because in some cases it is difficult to do so, but you should try describing it when you can.

Ideally, when writing the annotation for each step in the drawing, imagine that you are teaching a kid how to draw. How would you describe to them about the object you are drawing?

Figure 3.4: Final version of the requirements. The [Bad Example](#) links to counter-examples of the requirements, and [Good Example](#) links to good examples. When turkers click on the links, they are directed to the examples illustrating the corresponding requirement. This design helps them to understand the requirements better and provides high-quality annotations

turkers to switch between questions; originally, we displayed all questions on one page, but some found it time-consuming to scroll from the later questions back up to the instructions, so we decided to display one question at a time. We show one question from the final qualification in Figure 3.7. Each question is a mock-up of the main task interface, and turkers need to determine whether every step satisfies all the requirements. We also include hints on which requirement the question is testing to encourage turkers to revisit the requirements and form better understanding of the task. To see the full test, refer to: https://erinzhang1998.github.io/portfolio/amazon_qual.

3.1.3 Results

In our first pilot, we used prompts in the forms of *adjective* × *noun*. The list of adjectives includes: *happy*, *sad*, *surprised*, *sleepy*, *love-struck*, *evil*; the list of nouns includes: *person*, *kid*, *cat*, *bear*, *dog*, *sheep*, *jellyfish*, *cup of boba*, *apple*, *burger*, *sun*, *moon*, *star*. We want to see what sketches and text descriptions annotators would provide for prompts that ask for imaginative beings not in this world and include novel compositions of unrelated concepts, such as *evil apple* or *love-struck moon*. With

these creative prompts, we hope to collect data that contain interesting compositions of the same geometric shapes and descriptions across different objects. We can then learn models that can, for example, generate circles to be different parts in different objects: eyes, moon, cherries, and angel halo.

Since we only collected 55 sketches, we were able to manually examine every sketch, and we found many creative sketches, such as a subset of the sketches that were collected for prompts with the word *sleepy*, shown in Figure 3.8. From the pilot, we have learned that sketching was a good domain to study machine creativity, because examining the sketches in Figure 3.8, we saw that everybody had their own ways of illustrating *sleepy*: some people drew the letter *z* to symbolize sleeping; some drew a speech bubble with the word *yawn*; some drew sheep, cat, and dog lying down. However, one issue was that turkers took a long time, on average 30 minutes, to complete one sketch and provide descriptions, and we did not have the resources to spend such long time on a single sketch.

The second problem was that it was difficult to understand how some annotators interpreted the prompts through their sketches. We give 4 examples in Figure 3.9. Indeed, sketching is by its nature very subjective, a common challenge in creative AI.

The third problem, the most concerning one, was that the part descriptions did not align well with the sketches: some annotators failed to describe every part they drew in a step, or they described parts not in the annotated step. An example is shown in Figure 3.10.

3.2 Contrasting Sketch Text Dataset

3.2.1 Overview

In response to the pilot results, we reconsider the data collection pipeline. We examined existing sketch datasets to see how their annotations could facilitate our data collection (dataset details are explained in Section 2). To shorten the time spent on sketching, we no longer asked annotators to sketch and instead only asked them to describe sketches in the QuickDraw dataset. Although we were no longer able to design text prompts ourselves and collect creative sketches like the ones in the previous pilot, we solved the problem that it was difficult to understand how some sketches illustrate the given prompts.

The most important objective that this dataset should fulfill is allowing us to study how sketches

share similar semantic parts and descriptions that are adapted to be used for different objects. Therefore, we must resolve the issue that some text descriptions we collected in the pilot either did not describe everything drawn in a step or described more than what was drawn. Therefore, we used semantic part annotations from the Sketch Perceptual Grouping (SPG) dataset, which provides semantic part labels for each stroke in 20,000 sketches from the QuickDraw dataset. In this way, annotators did not need to spend time thinking about how to segment sketches into parts themselves.

Moreover, to help annotators coming up with creative ways to describe the semantic parts, in each task, we presented a pair of sketches with contrasting features, implicitly priming them to describe the visual differences.

3.2.2 Interface Design

Main Task Interface

When collecting the previous prompt-guided dataset, we relied on testing the interface with students in the lab to determine our design, but we observed performance differences between the students and turkers, such as amount of variety in the sketches, amount of time spent on the task, and common confusions related to understanding the requirements. Therefore, when collecting the contrasting sketch text dataset, we deployed several pilots on AMT to design the new interface. In Figure 3.11, we show how the main task interface progressed from the first pilot to the final version used to collect the entire dataset.

To better study how similar words are used differently across sketches, we changed to collect adjective phrases (Figure 3.11b, 3.11c) from collecting whole sentences (Figure 3.11a). We juxtaposed two sketches and highlighted the parts to be annotated in different colors to help annotators notice the contrasting features. Moreover, this design expedited the annotation process, since it was easier for people to perform contrasting tasks than to generate descriptions from a single sketch.

Instruction and Requirement

At the beginning, the instruction limited the annotators to provide three types of descriptions: shape, size, and position. However, in order to collect creative descriptions, we lifted restrictions on the type of words and only required annotators to fill in the blank with adjective phrases. We also provided some examples of adjective phrases in common sentences, unrelated to our task, for

annotators to better understand their usage (Figure 3.12).

Since we simplified the HIT from 3 sub-tasks, sketching for the prompt, segmenting the sketches, and describing each step, to only asking for part descriptions, the requirements are much easier to write. We received less feedback from the annotators on being confused about the kind of sketches we wanted and the definition of semantic parts, which are now automatically highlighted in the sketches.

We relied on the examples in the instruction to give annotators an idea of what descriptions we wanted. Some examples that we used in the tasks are shown in Figure 3.14. However, the downside was that, primed by the examples, annotators described the parts using words in the examples instead of coming up with a variety of descriptors, and we observed this behavior in the pilots. Therefore, we emphasized an additional requirement that asked the annotators to not limit themselves to words in the examples, and they should use any words that could illustrate the parts well.

Since in the future we want to use our dataset to learn models that can generate sketch parts from text descriptions, the text should pertain to visual properties of the parts, so we required that *Do not use adjectives that fail to describe specific visual properties of the objects in the sketches*. A caveat was that some annotators might consider descriptions about emotions expressed in the sketches unrelated to visual properties, since they are abstract compared to words like *rectangular* and *large*. We wanted to collect annotations for face sketches, and parts like *eyes* and *mouth* can be *smiley* or *sad*, so we added that adjectives describing emotions were allowed.

The requirements used in the final version is shown in Figure 3.13.

Qualification

We prepared 10 qualification questions, all yes/no questions. We used the qualification test to train turkers to understand the requirements better. Each question had a hint that stated which requirement and examples were helpful for solving the questions. The purpose of the qualification test was not to trick annotators but to ensure speed and quality of the annotation. We show one question from the qualification in Figure 3.15. To see the full test, refer to: <https://erinzhang1998.github.io/portfolio/v2qual>. At the end, there are 88 annotators who worked on our tasks.

3.3 Dataset Summary

Our dataset contains sketches, their semantic part annotations, and descriptions for every part in a sketch. The sketches comes from the QuickDraw dataset (Ha & Eck, 2017), and the semantic part annotations come from the SPG dataset (Li et al., 2018); both datasets are explained in details in Section 2.

We annotated for 2 categories of sketches: face and angel. For angel sketches, we annotate for the parts *halo*, *eyes*, *nose*, *mouth*, *body*, *outline of face*, and *wings*. For face sketches, we annotate for the parts *eyes*, *nose*, *mouth*, *hair*, *outline of face*.

	Face	Angel
Number of contrasting pairs	2515	3060
Number of distinct words	833	1107
Number of sketches	572	787

Table 3.1: Dataset statistics by category.

	Face					Angel						
	eyes	nose	mouth	hair	face	halo	eyes	nose	mouth	face	body	wings
Number of sketches	334	572	572	104	572	558	114	8	80	732	781	779
Number of distinct words	228	360	325	152	314	365	112	21	88	379	425	534
Number of contrasting pairs	689	401	687	126	612	559	114	8	80	733	785	781

Table 3.2: Dataset statistics by sketch parts. The phrase *contrasting pair* refers to a pair of sketches with contrasting features that are presented to the annotators for descriptions of their parts.

In Table 3.2, we present dataset statistics broken down by semantic parts, and, in Table 3.1, we show the same statistics by sketch category. In Figure 3.16, we show 100 most frequently used words in our dataset.

The *Creative Birds* and *Creative Creatures* datasets collected by DoodlerGAN (Ge et al., 2020) contain 2 categories, like ours, but there are 9k sketches in each category, and ours contains one tenth as many. Although we fall short on the number of sketches and the variety of sketching styles, we approach creativity from a completely different angle: we focus on how people compose similar basic shapes to create sketches of different categories and adapt similar language to describe different sketch parts, as explained in Section 1 and Figure 1.1.

Our face sketches contain 5 different parts, and angel sketches has 7 parts; there are 7 parts in Creative Birds and 16 parts in Creative Creatures. Again, we do not annotate for as many parts as

DoodlerGAN if compared directly with their dataset. However, in our dataset, every part in every sketch has at least 2 different language annotations describing the visual features, while DoodlerGAN has no language annotation. With text descriptions, in addition to part-based generative model, we can study how to generate similar shapes from different text descriptions. In Figure 3.17, we show how people describe *circles* using various descriptors (the figure only contains a subset of ways describing circles). In Figure 3.18a, we show the variety of language in describing a simple smiley face. Similarly, an example of creative description of different parts in angel sketches is shown in Figure 3.18b.

	Face					Angel						
	eyes	nose	mouth	hair	face	halo	eyes	nose	mouth	face	body	wings
small	192	143	65	2	9	150	53	8	25	202	64	47
oval	123	7	11		207	258	1			145	42	23
circular	164	12	11	2	192	80	14	1	2	276	10	11
round	175	18	2	1	130	91	8			268	22	45
wide	94	16	193		74	84	2		12	50	120	111
open	118	7	106		37	54	6		1	79	137	26
large	98	41	58		22	66	6			131	54	85
curve	23	52	221	3	17	18	9	1	27	10	58	101
triangular	9	20	9	2	11	10				14	294	42
line	92	27	56	19	1	15	22	3	14	3	93	9

Table 3.3: This table shows how many times the top-10 most frequently used word in the dataset is used to describe different parts in face and angel sketches. For example, *small* is used in 192 descriptions for eyes in face sketches, and it is used 150 times to describe angel halos. People use the same words differently depending on the sketching context.

Moreover, using this dataset, we can study how people compose concepts (e.g. *large+round*) and apply the same concepts to different semantic parts (e.g. *large round eyes* and *large round halo*). We observe that face and angel sketches share 486 words in total, and for the 200 words that occur most frequently in the dataset (all used in at least 13 part descriptions), 182 are shared across the two categories. Therefore, we believe that the datasets would allow us to study how people use similar words when sketching different objects. Although we cannot conclude in this project, we have already found some examples of the meaning of descriptors varied across sketches, after examining a few words. Figure 3.19 shows an example. In this example, *spiky* is used to mean either short hair in face sketches or pointy, edgy wings in angel sketches. Moreover, there is an implicit orientation change between the two categories of sketches: while *spiky hair* implies “spike” in the vertical direction, the phrase, *spiky wings*, implies jigsaws pointing out from the angel body, horizontally.

Bad Example 2

Step 1

Prompt	Canvas	Annotation
A house		<input type="text" value="Annotate an item..."/> <button>Add</button> triangle <button>Delete</button>

This step violates: Name the object in the annotation.
A roof is drawn, so the annotation should include the name, "roof". A good annotation of this step would be "triangular roof".

Step 2

Prompt	Canvas	Annotation
A house		<input type="text" value="Annotate an item..."/> <button>Add</button> triangle <button>Delete</button> rectangle <button>Delete</button>

This step violates: Name the object in the annotation.
The body of the house is drawn, so the annotation should include the name, "body of the house". A good annotation of this step would be "rectangular body of the house".

Step 3

Prompt	Canvas	Annotation
A house		<input type="text" value="Annotate an item..."/> <button>Add</button> triangle <button>Delete</button> rectangle <button>Delete</button> a cross in a rectangle as the left window <button>Delete</button>

This step violates: The word "left" always refers to this side: ←
Annotation used the word "left", but a window is drawn on the right.

Step 4

Prompt	Canvas	Annotation
A house		<input type="text" value="Annotate an item..."/> <button>Add</button> triangle <button>Delete</button> rectangle <button>Delete</button> a cross in a rectangle as the left window <button>Delete</button> a cross in a rectangle as the right window <button>Delete</button>

This step violates: The word "right" always refers to this side: →
Annotation used the word "right", but a window is drawn on the left.

Figure 3.5: An example used to explain the requirements to turkers.

Qualification Test

Each question is an example of how someone, given a prompt, would draw and then annotate for the object at each step. You will need to determine if a step satisfies all the must satisfied requirements by clicking "Yes" or "No".

Note that:

- Strokes drawn in the current step are shown in black while strokes drawn in previous steps are shown in gray
- Annotation for the current step has border around the box

For a question to be counted as correctly answered, you will need to correctly determine the validity of **all** steps.

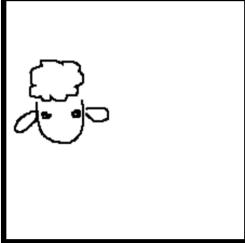
[Question 1](#) [Question 2](#) [Question 3](#) [Question 4](#) [Question 5](#) [Question 6](#) [Question 7](#) [Question 8](#) [Question 9](#) [Question 10](#)

Figure 3.6: Navigation bar in the qualification test.

Step 1

Q: Does this step satisfy all the requirements? Yes No

Hint: this step is intended to test your understanding of Requirement 1. Check [Bad Example 1](#).

Prompt	Canvas	Annotation
sheep		<input type="text" value="Annotate an item..."/> Add <div style="border: 1px solid blue; padding: 2px;">sheep head</div> Delete

Step 2

Q: Does this step satisfy all the requirements? Yes No

Prompt	Canvas	Annotation
sheep		<input type="text" value="Annotate an item..."/> Add <div style="border: 1px solid blue; padding: 2px;">sheep head</div> Delete <div style="border: 1px solid blue; padding: 2px;">sheep body that looks like fluffy cloud</div> Delete

Figure 3.7: An example of the questions in the qualification test.



Figure 3.8: Sketches collected for prompts with the descriptor *sleepy* in the prompt-guided sketch text dataset. People were very creative and came up with various ways to illustrate the concept *sleepy* in their sketches.

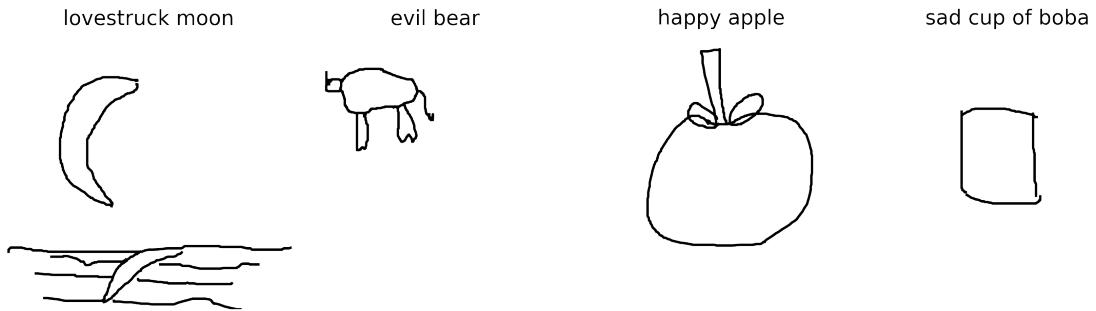


Figure 3.9: A challenge studying creativity in humans: sketching is very subjective, and it is difficult to determine how some sketches illustrate the prompt and judge their annotation quality. This figure shows 4 examples from the prompt-guided sketch text dataset that are hard to interpret.

(Step 1) circle as the face of the cat (Step 2) oval as the body of the cat (Step 3) oval as the legs of the cat



(Step 4) stripes as the mustache of the cat



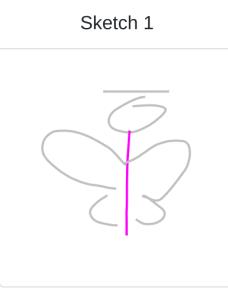
Figure 3.10: A major issue with the prompt-guided sketch text dataset: many text descriptions do not align with the drawings in each step. In this example, annotations for step 1 and 2 are correct, but in step 3, although the description says the sketch contains *legs of the cat*, the person actually drew the cat face, ears, whiskers, and legs. Moreover, in step 4, the annotator annotated for parts drawn in the last step. Resolving this problem was very important because we needed semantic part annotations and their corresponding text descriptions to study abstract concept composition.

Sketch Category	Sketches
angel	<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>Sketch 1</p>  </div> <div style="text-align: center;"> <p>Sketch 2</p>  </div> </div> <p>Describe differences between the angel bodies (strokes in magenta color) in the two sketches.</p>

(a) In the first pilot, we ask annotators to describe differences between the two sketches in full sentence.

Sketch Category	Sketches
angel	<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>Sketch 1</p>  </div> <div style="text-align: center;"> <p>Sketch 2</p>  </div> </div> <p>Q1: Compared to Sketch 2, Sketch 1 draws [] angel wings (strokes drawn in olive color).</p> <p>Q2: Compared to Sketch 1, Sketch 2 draws [] angel wings (strokes drawn in olive color).</p> <p><input type="checkbox"/> If someone is shown the two sketches, the person can pick out one sketch based on the provided differences.</p>

(b) Compared to the previous pilot (3.11a), we still ask explicitly the differences between the sketches but limit annotations to adjective phrases.

Sketch Category	Sketches
angel	<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>Sketch 1</p>  </div> <div style="text-align: center;"> <p>Sketch 2</p>  </div> </div>
	<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;">  </div> <div style="text-align: center;">  </div> </div>
	<p>body (object drawn in magenta color).</p> <p>body (object drawn in magenta color).</p>

The phrases satisfy all requirements (shown on the left).

(c) Similar to the last pilot (3.11b), the final version asks for adjective phrases, but we do not state explicitly that annotators should describe the visual differences to allow for more creativity.

Figure 3.11: Different versions of the main task interface in chronological order. Final version used to collect the entire dataset is 3.11c.

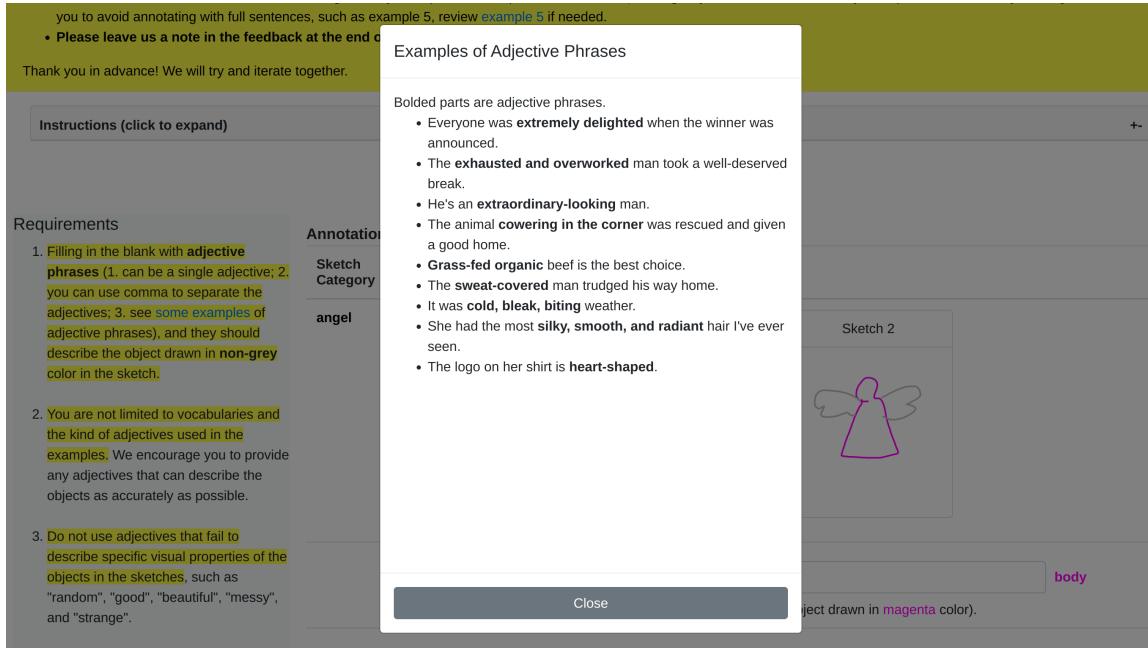


Figure 3.12: A pop-up window containing examples of adjective phrases, used in collecting contrasting sketch text dataset in Section 3.2. It can be opened from the side panel by clicking on [some examples](#). Note that we used examples that are unrelated to our task on purpose, because annotators tended to repeat words in examples, so we tried to not bias them.

Requirements

1. Filling in the blank with **adjective phrases** (1. can be a single adjective; 2. you can use comma to separate the adjectives; 3. see [some examples](#) of adjective phrases), and they should describe the object drawn in **non-grey** color in the sketch.
2. You are not limited to vocabularies and the kind of adjectives used in the examples. We encourage you to provide any adjectives that can describe the objects as accurately as possible.
3. Do not use adjectives that fail to describe specific visual properties of the objects in the sketches, such as "random", "good", "beautiful", "messy", and "strange".
4. Do not use comparative and superlative adjectives, such as "wider" and "widest" (use "wide" instead).

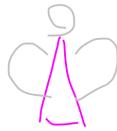
Figure 3.13: Requirements used in collecting the contrasting sketch text dataset (Section 3.2).

Example 2 (Bad)

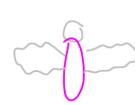
Sketch
Category Sketches

angel

Sketch 1



Sketch 2



circular

angel body
(object drawn in magenta color).

rectangular

angel body
(object drawn in magenta color).

Problems:

The adjective, "circular", **does not** describe the angel body in Sketch 1: "triangular" would be more suitable.

The adjective, "rectangular", **does not** describe the angel body in Sketch 2: "oval" would be more suitable.

Using "circular" and "rectangular" are too far-fetched and are considered incorrect annotations. **Check Requirement 1.**

One correct annotation:

Sketch 1: "triangular" angel body
Sketch 2: "oval" angel body

Example 4 (Good)

Sketch
Category Sketches

angel

Sketch 1



Sketch 2



thin rectangular

angel body
(object drawn in magenta color).

wide rectangular

angel body
(object drawn in magenta color).

Satisfy all requirements.

You are not limited to adjectives used in the examples.

Requirement 2 encourages you to use any adjectives that can reasonably describe the objects.

One way to come up with adjectives is to think about what adjectives are necessary to distinguish the objects in the two sketches.

Figure 3.14: 2 examples used in the instructions to collect the contrasting sketch text dataset (Section 3.2).

Question 3

Hint: review requirement 1 and [example 6](#).

Sketch Category	Sketches
pig	<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>Sketch 1</p>  </div> <div style="text-align: center;"> <p>Sketch 2</p>  </div> </div> <div style="display: flex; justify-content: space-around; margin-top: 10px;"> big circular pig ears small curly pig ears </div> <div style="display: flex; justify-content: space-around; margin-top: 5px;"> (object drawn in red color). (object drawn in red color). </div>

Q: Do the answers above satisfy all the requirements? Yes No

Figure 3.15: Question 3 from the qualification test used to collect the contrasting sketch text dataset (Section 3.2). We show a hint at the beginning of each question telling the annotators which requirement this question is testing. In this way, we encourage them to review the requirements so that they have a good understanding of the task and can provide high-quality annotations in the real HIT. The question interface is the same as main task interface that annotators will see when they annotate. The 1-to-1 mock-up helps them to be familiar with the workflow.

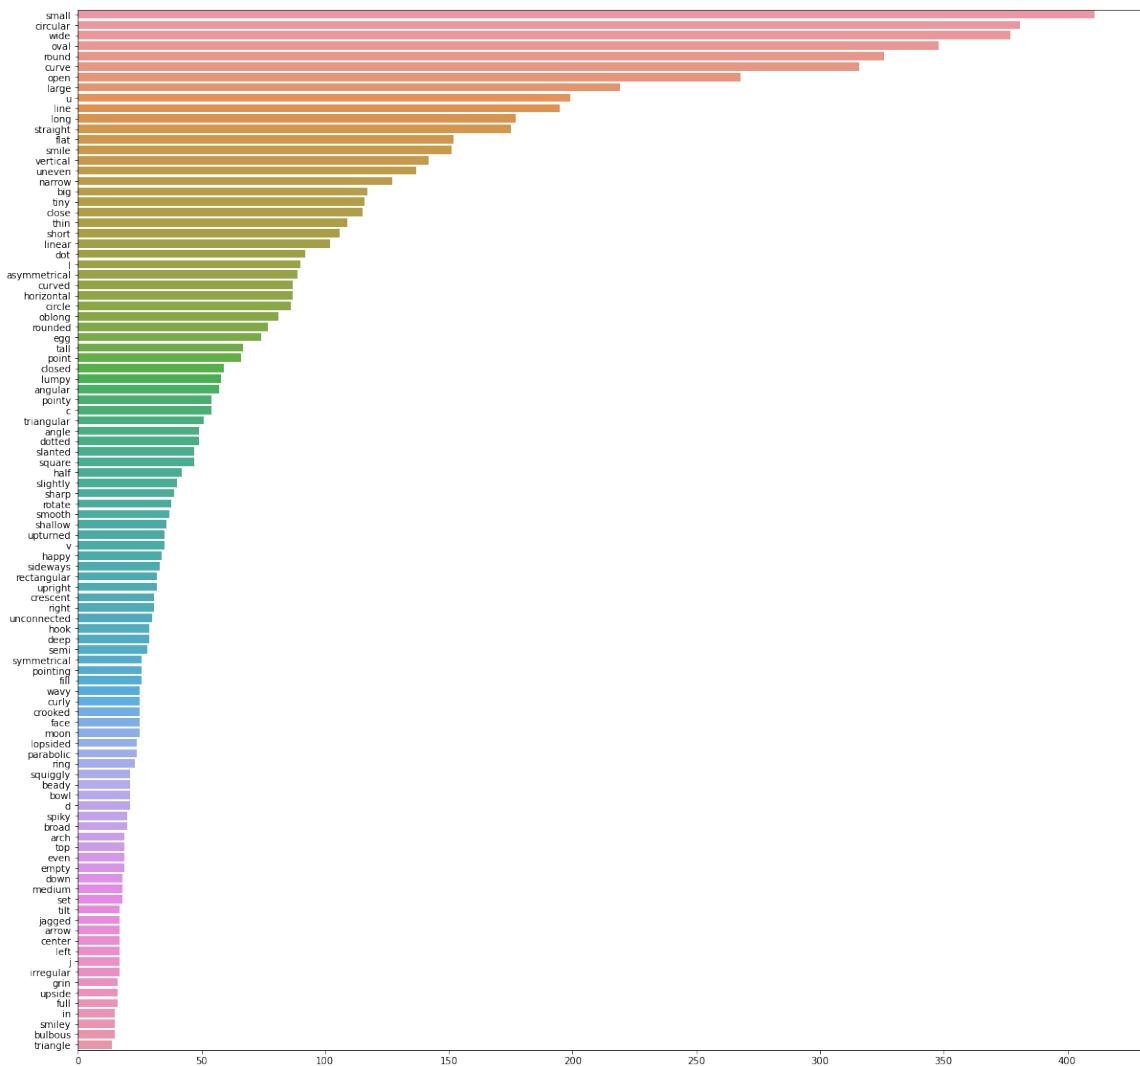


Figure 3.16: Top 100 most frequent words in the dataset corpus.

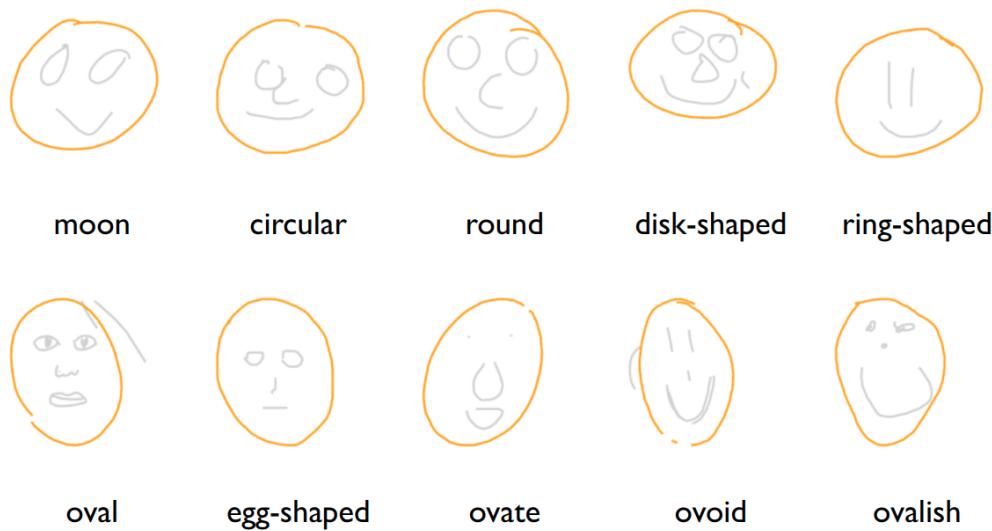


Figure 3.17: .

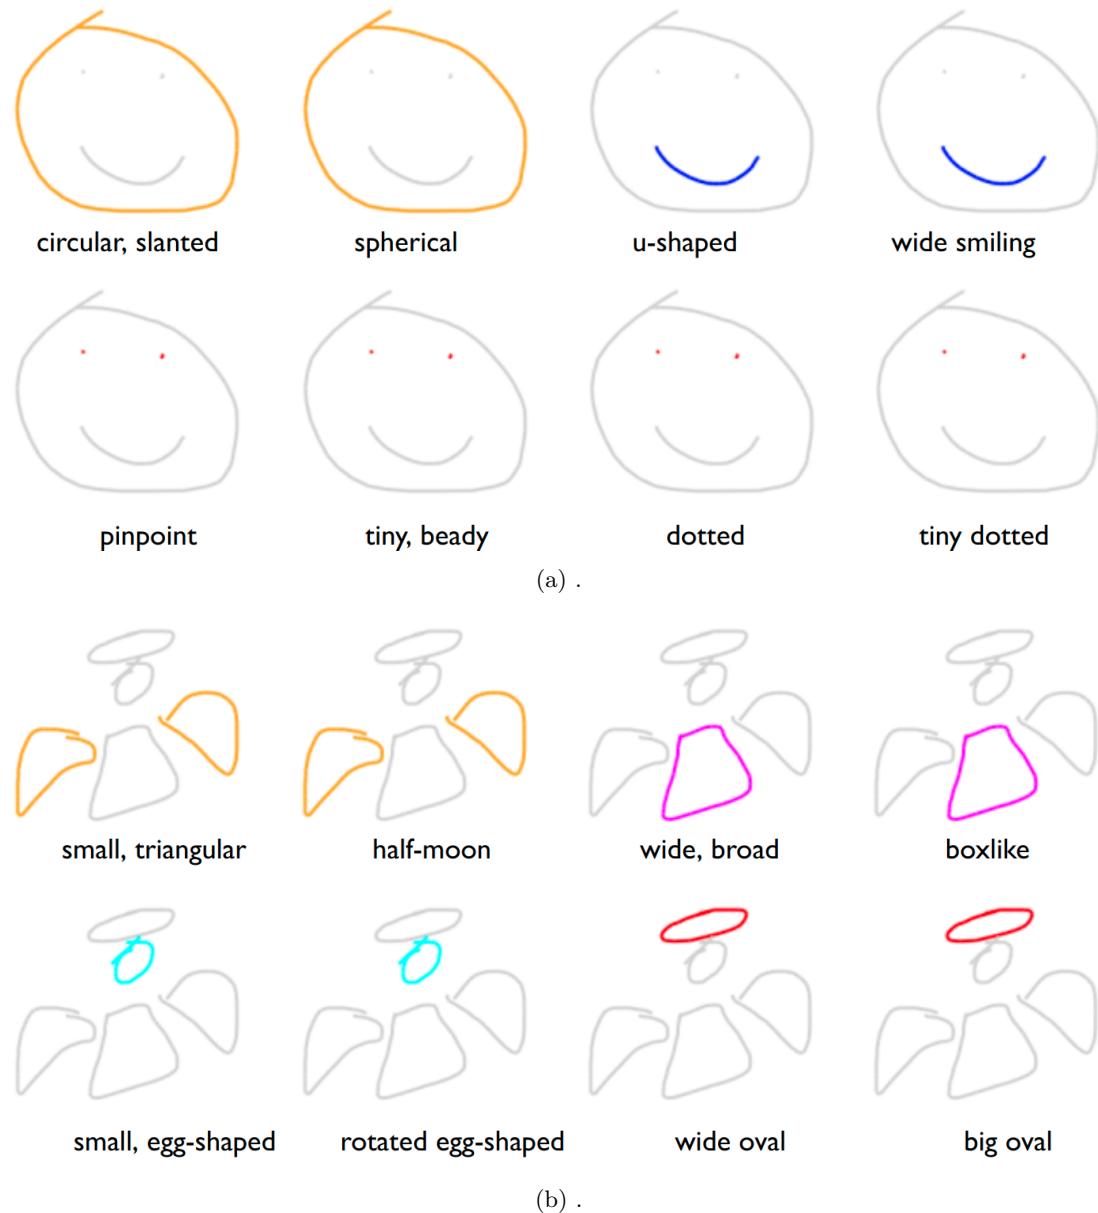


Figure 3.18: .



Figure 3.19: Left: face sketches whose hair descriptions contain the word *spiky*. In faces, the word refers to short and coarse texture of hair. Right: angel sketches whose wing description contains the word *spiky*, which is used to represent the pointy edges on wings.

Chapter 4

Modeling

4.1 Task Definition



(a) t_1 : wide triangular body



(b) t_2 : small rectangular body

Figure 4.1: Two angel sketches, s_1 on the left and s_2 on the right, and their part annotations, t_1 and t_2 . The task is for CLIP to determine which sketch matches a given t_j .

Given two sketches (s_1, s_2) and their part annotations (t_1, t_2), such as the pair shown in Figure 4.1, the task is to determine which sketch t_j should be paired with. We use this task to evaluate the joint vision-language embedding space of CLIP, which is often used as part of the objective function for models that generate image from text; the objective function often involves maximizing the cosine similarity between the generated image and the provided text (or some variants engineered to work for the particular latent space of the generator used) (Frans et al., 2021; Patashnik et al., 2021; Gal et al., 2021; Ramesh et al., 2022). Through this task, we can study how well CLIP can recognize different visual concepts in sketches, and the experiments can help us determine if CLIP can be used

in similar ways to guide part-based sketch generation from language. Moreover, since we give CLIP the same pairs that were given to the annotators, we can learn if CLIP can understand how people are using language to describe the visual features of the sketches.

4.2 Method

CLIP stands for Contrastive Language-Image Pre-training, and it seeks to learn image representations that transfer to a wide range of downstream tasks such as image classification on datasets with a wide variety of source domains. To do so, it uses the task of pairing images with their corresponding captions for pre-training. CLIP has two main parts: a text encoder and an image encoder, and both can be transformers. To model our dataset, we fine-tune the pre-trained ViT-B/32 CLIP model using the Python `clip` package.

In this section, we will first introduce the details of the CLIP contrastive objective (Section 4.2.1). We will then introduce the pre-trained transformer-based text and image encoders in Section 4.2.2 and 4.2.3, respectively. Lastly, in Section 4.2.4, we explain how we preprocess sketches and texts for fine-tuning and how we use CLIP to perform the task in Section 4.1.

4.2.1 Contrastive Objective

Pre-training on an enormous amount of image-caption data with a contrastive objective, CLIP is able to learn robust vision-language joint embedding that allows it to perform on par with state-of-the-art (SOTA) models learnt with supervised objectives (Radford et al., 2021). During pre-training, for a batch of N (text, image) pairs, CLIP obtains N image features, I_1, \dots, I_N , and N text features, T_1, \dots, T_N , from the encoders. It then calculates dot-product between each pair of normalized vectors, where there are $N \times N$ possible pairing in total, as shown in the grid in Figure 4.2. This grid is a logit matrix X of dimension $N \times N$ and $X_{ij} = I_i \cdot T_j$. If we normalize the i -th row ($i \in [N]$) through softmax ($\frac{\exp\{X_{i,i}\}}{\sum_{j=1}^N \exp\{X_{i,j}\}}$), we obtain a distribution over the captions representing the likelihood of pairing the j -th caption with the i -th image. Similarly, normalizing the j -th column through softmax gives a distribution over all the images of how likely an image pairs with the j -th caption.

As explained in Radford et al. (2021), we can treat each batch as containing N visual concepts expressed through language. By normalizing the rows and columns into distributions, we perform

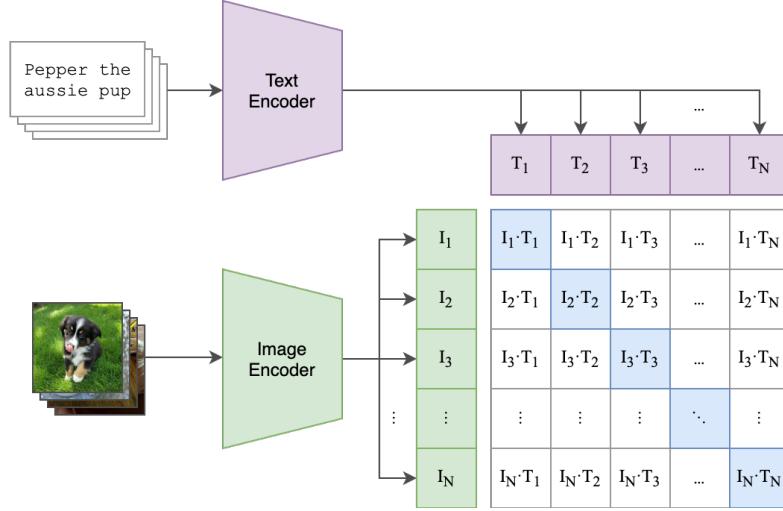


Figure 4.2: CLIP uses contrastive instead of generative objective for pre-training to learn joint vision-language embeddings. This image is taken from Figure 1 in the original CLIP paper (Radford et al., 2021). Each grid contains the dot-product of normalized image, text features.

classification over the N captions for the images and classification over the N images for the captions. Moreover, we know that the ground-truth is pairing the i -th image with the i -th caption. Therefore, for classification over captions, we have the ground-truth vector Y_I ; for classification over images, we have the ground-truth vector Y_T , and we know:

$$Y_I = Y_T = [1 \quad 2 \quad \dots \quad N]^T$$

Similar to standard multi-class classification, we can use cross-entropy loss as our objective function. The loss $L_I(X, Y)$ of selecting the correct caption for each image:

$$L_I(X, Y) = \frac{1}{N} \sum_{i=1}^N -\log \frac{\exp\{X_{i,i}\}}{\sum_{j=1}^N \exp\{X_{i,j}\}} \quad (4.1)$$

The loss $L_T(X, Y)$ of selecting the correct image for each caption:

$$L_T(X, Y) = \frac{1}{N} \sum_{j=1}^N -\log \frac{\exp\{X_{j,j}\}}{\sum_{i=1}^N \exp\{X_{i,j}\}} \quad (4.2)$$

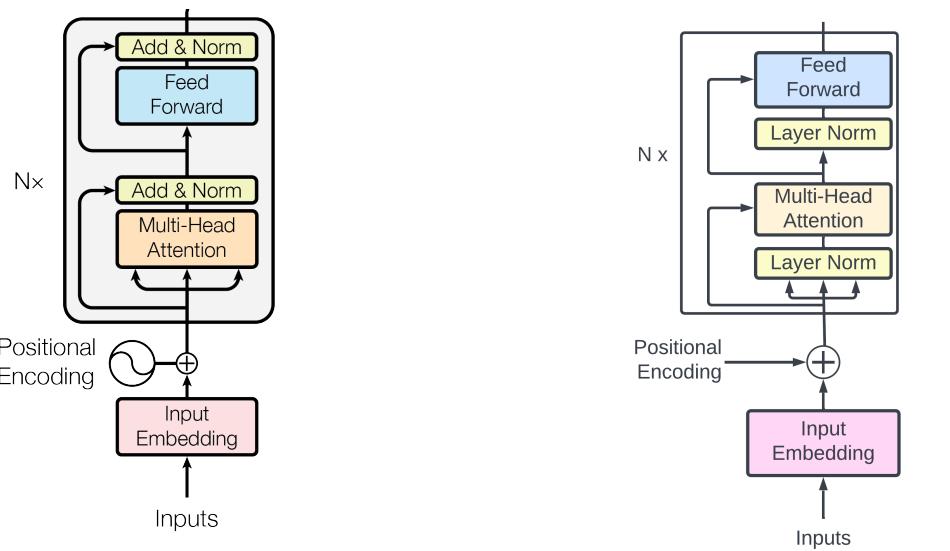
The final loss is defined as:

$$L = \frac{1}{2}(L_I(X, Y) + L_T(X, Y)) \quad (4.3)$$

In order to minimize the cross-entropy loss, the model has to increase the logits on the diagonal of X , which means that it has to learn an embedding space where the feature vectors of the i -th image and the i -th caption are similar and the feature vectors of the other pairs are far apart. The term *contrastive* in the title of CLIP comes from this objective to pull together the real pair in the joint embedding space.

4.2.2 Text Encoder

Overview of Transformer Architecture



(a) Encoder architecture used in original Transformer paper; the figure is taken from Figure 1 in Vaswani et al. (2017).

(b) Encoder architecture of CLIP (Radford et al., 2021).

Figure 4.3: Text encoder architecture: the text encoder is a transformer with stacked layers of self-attention and feed-forward network sublayers. The main implementation difference between the original Transformer (on the left) and the transformer implemented by CLIP is the placement of layer normalization.

The text encoder of CLIP is based on the transformer architecture introduced in Vaswani et al. (2017). Transformer relies only on self-attention mechanism to compute a representation for the input sequence. In this way, it alleviates the computation inefficiency and difficulty capturing long-range dependencies witnessed in recurrent layers.

Figure 4.3a and 4.4 are the same figures used in Vaswani et al. (2017) to illustrate the transformer

architecture. In Figure 4.3a, Vaswani et al. (2017) gives an overview of the encoder architecture. Firstly, tokenized texts go through an embedding layer; the input embeddings are summed with learned position embeddings that inject information on order of the sequence. The input is computed using Byte Pair Encoding (BPE) with a 49,152 vocabulary. As explained in the GPT-2 paper, BPE, a sub-word tokenization scheme, strikes a good balance between word-level and character-level word embeddings, since one works well with common words and the other with rare sequences (Radford et al., 2019).

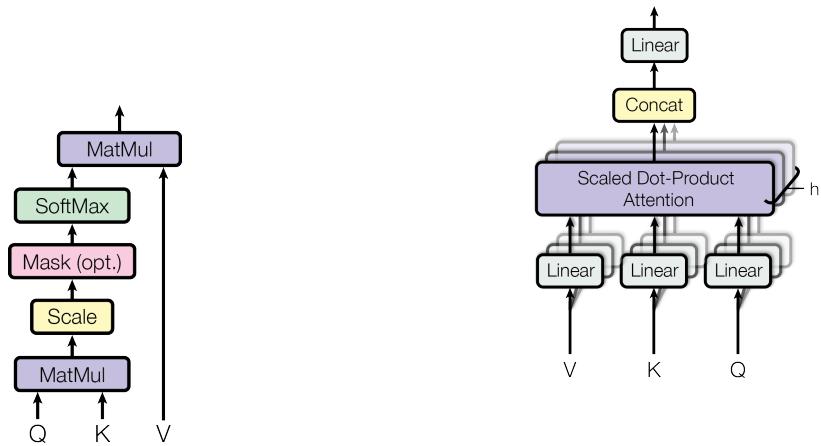


Figure 4.4: An illustration of the scaled dot-product attention on the left; on the right, an illustration of the multi-head self-attention mechanism (MSA) used in every layer of the transformer. Both figures are from the original Transformer paper (Vaswani et al., 2017).

Next, the input is passed through stacked layers multi-head self-attention (MSA) mechanism followed by point-wise feed-forward networks (FFN). In the version of CLIP that we use, the text encoder is a 12-layer transformer, and the model dimension is 512, $d_{model} = 512$, meaning that output of the initial embedding layers, MSA, FFN all have dimension 512.

Compare to the formulation $\text{LayerNorm}(x + \text{Sublayer}(x))$ used in Vaswani et al. (2017) (Figure 4.3a), the CLIP text encoder uses $x + \text{Sublayer}(\text{LayerNorm}(x))$ (Figure 4.3b), where Sublayer refers to either MSA or FFN. Each layer still contains residual connection and layer normalization, but the order is switched between Vaswani et al. (2017) and Radford et al. (2021).

Multi-Head Self-Attention

As illustrated in Figure 4.4, given query, key, value matrices Q, K, V (in our case, all three matrices equal to the input text embeddings), transformer uses different linear projections to create

multi-head attention, and Vaswani et al. (2017) explains the benefit of multi-head attention as allowing the model to attend simultaneously to multiple representation subspaces of the input.

$$\begin{aligned} \text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \\ \text{head}_i &= \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \\ \text{Attention}(Q, K, V) &= \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \end{aligned} \quad (4.4)$$

ViT-B/32 uses a version with $h = 8$ attention heads, so $W_i^Q, W_i^K \in \mathbb{R}^{d_{model} \times d_k}$ and $W_i^V \in \mathbb{R}^{d_{model} \times d_v}$, where $d_k = d_v = \frac{d_{model}}{h} = \frac{512}{8} = 64$. At the end of the multi-head attention mechanism, the weighted combination of values from each head is concatenated together and passed through a linear layer, represented here as $W^O \in \mathbb{R}^{(d_v \times h) \times d_v}$. CLIP uses the “Scaled Dot-Product Attention” in Vaswani et al. (2017), illustrated in details in Figure 4.4. The dot products between query and key determine the weights that are used to sum the values; in this way, we have a contextualized representation; compared to convolutions that use static kernels, attention weights are dynamic.

4.2.3 Image Encoder

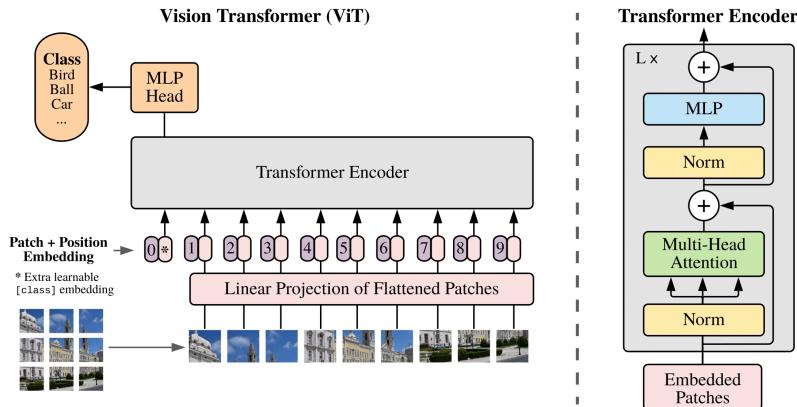


Figure 4.5: Vision Transformer (ViT) architecture (Dosovitskiy et al., 2020).

The image encoder of CLIP uses Vision Transformer (ViT) introduced in Dosovitskiy et al. (2020). The architecture of ViT is based on the original transformer introduced in Vaswani et al. (2017). In order to reuse the transformer model, ViT needs to first turn an image of size $H \times W \times C$, (H, W, C stands for image height, width, channel size, respectively), into a sequence of “tokens”,

similar to the text input. To do so, Dosovitskiy et al. (2020) reshapes the image to size $N \times (P^2 \cdot C)$, where N is the number of patches and P the patch size; the reshaped image can be seen as a sequence of N image tokens, each having a dimension of $P^2 \cdot C$. Each image token is then passed through a linear layer to be mapped to dimension D , similar to the model dimension d_{model} earlier.

In the version of CLIP that we used, all input images have dimension 224×224 , the patch size $P = 32$, and the output of the initial embedding layers, MSA, FFN all have dimension 768 ($D = 768$). As explained in Dosovitskiy et al. (2020), before passing into the transformer, we also need to prepend a `[class]` token at the front the sequence, whose embedding at the last layer of the transformer will be used as the representation for the entire image. In this way, each image is represented as a sequence of $7 \times 7 + 1 = 50$ tokens, illustrated as purple boxes in Figure 4.5. The pink boxes that are right next to the patch embeddings represent position embeddings, with a similar function to encoder sequence order information as in the text transformer. For the CLIP image encoder, layer normalization is applied to the input before passing into the transformer and to the output at the last layer.

4.2.4 Fine-Tuning CLIP

We fine-tune CLIP with our dataset containing (sketch, part description) pairs. Our sketches come from the QuickDraw dataset (Ha & Eck, 2017), and the semantic part annotations come from the SPG dataset (Li et al., 2018). We collected text descriptions from turkers for every part in face and angel sketches; details of the data collection process is in Section 3.2, and we summarize the dataset in Section 3.3.

Sketch Preprocessing

The QuickDraw sketches are stored in vector format: each sketch is composed of a sequence of n strokes $S_i, i \in [n]$, and each stroke S_i is a sequence of vectors $(\delta x, \delta y, p, l)$. δx and δy are changes in the x, y coordinates with respect to the previous point; for the first point, its coordinate is with respect to the point (25, 25). All points are assumed to be drawn on a 256×256 canvas. $p = 1$ if the point is the last point in the current stroke, and $p = 0$ otherwise. The SPG dataset provides annotation for semantic segmentation of the sketches, and l is an integer representing which semantic part the current point belongs to. For face sketches, the parts can be *eyes*, *nose*, *mouth*, *hair*, or *outline of face*; for angel sketches, the parts can be *halo*, *eyes*, *nose*, *mouth*, *body*, *outline of face*, or *wings*.

During training, we render the sketches from the vector format into PNG images using the python package `cairocffi`. To augment the data, we did: (1) sample stroke width uniformly from the range [2.5, 12.5]; (2) rotate the sketches randomly by sampling an angle from the range $[-\frac{\pi}{4}, \frac{\pi}{4}]$; (3) scale the sketches with the scale factor sampled from the range [0.75, 1.25]; (4) translate the strokes vertically or horizontally with 0.15 as the maximum absolute fraction translation.

Text Preprocessing

We used the `spacy` package to preprocess the text. `spacy` provides trained natural language processing pipeline and includes models for, for example, token-to-vector and part-of-speech tagging. We use the `en_core_web_sm` pipeline and its lemmatizer to reduce words to their basic forms. Moreover, we lower-case all words and remove punctuation, a list of which is provided by Python `string` package, `string.punctuation`. We also remove words like *shaped*, *sized*, *and*, *like*, since they act like stop words and do not provide additional visual descriptions of the sketches. Text descriptions are also tokenized by CLIP’s tokenizer before passing into CLIP text encoder.

Fine-tuning Objective

We used the same loss function and model pipeline as the one used for pre-training CLIP, explained in details in Section 4.2.1. In this case, in each (image, text) pair, the *image* is a sketch, and the *text* is a description of one of the semantic parts in the sketch. Two examples are shown in Figure 4.1: through fine-tuning, CLIP should adjust its embedding space so that text feature for *wide triangular body* and image feature for the sketch in Figure 4.1a are close in cosine similarity.

Chapter 5

Results & Analysis

In this chapter, we present experiment results and analysis to understand the language used to describe sketch parts. We found the following:

1. CLIP, despite its large pre-trained corpus, cannot easily generalize to unseen category on the task of pairing sketches with their descriptions. This indicates that the relationship between natural images and language cannot be easily translated to sketches and their descriptions. (Section 5.1)
2. Average cosine distance has increased between a pair of descriptors that are used by annotators to differentiate two sketches, suggesting that word embeddings extracted from the fine-tuned CLIP can better reflect the contrasting nature of these pairs, compared to BERT or pre-trained CLIP. (Section 5.2)

5.1 Classification Experiment

As explained in Section 4.1, annotators provide descriptions (t_1, t_2) to differentiate the two sketches (s_1, s_2) presented to them, and from this process, we have the ground-truth pairing of (s_1, t_1) and (s_2, t_2) . We use CLIP to decide which sketch should be paired with description t_j , and the results are reported in Table 5.1.

Model	Face		Angel	
	Test	Dev	Test	Dev
zero-shot	53.8	54.6	56.4	57.2
finetuned on face	70.1	68.5	58.1	57.5
finetuned on angel	58.1	61.0	67.2	70.5
finetuned on face + angel	74.3	71.7	70.4	71.0

Table 5.1: Accuracy (%) of CLIP on choosing the correct sketch for a given text description from a pair of sketches. During annotation, the annotator is given the pair of sketches and provided descriptions of certain parts in the sketches for both sketches.

Compare to the zero-shot performance, fine-tuning on a single category would increase performance significantly on the category: for face, accuracy increased 16.3% on the test set after fine-tuning on face; performance increased 10.8% on the angel test set after fine-tuning on angel. The increase demonstrates that CLIP has adapted its embedding space and has learnt to associate sketches with their part descriptions. However, CLIP fine-tuned on a single category cannot generalize easily to unseen category: only 1.7% increase on angel when fine-tuned on face, and 4.3% increase on face when fine-tuned on angel. The two categories share vocabulary, and similar shapes appear in the two categories, but CLIP does not generalize as easily as expected. Concepts like *big* and *small* appear in both categories, but it is not guaranteed that CLIP can recognize them in sketches from unseen category.

Transferring from angel to face results in slightly better performance compared to transferring from face to angel (+4.3% vs. +1.7%). One likely explanation is that the angel category contains more sketches (787 vs. 572) and a larger vocabulary (1107 vs. 833), so CLIP performed better because it has seen more sketches and familiarized itself with the domain, and it has learnt from a wider variety of ways that sketches can be described. A second explanation is that angels contain all face parts (*eyes, nose, face outline, mouth*) except *hair*, but face sketches do not contain *body, wings*, and *halo*, which are the most definitive features of angels. Therefore, CLIP, after learning how face parts are described through a few angel sketches that contain them, can recognize face sketches better. However, when it encounters angel sketches with unseen parts and new usage of words seen in the face category, it fails to transfer the learnt concepts to angels.

Fine-tuning on both categories results in slightly better performance than fine-tuning on a single category. CLIP has the capacity to learn how different words are used in both categories, but it cannot transfer the learnt concept across categories. Moreover, as explained in Section 4.2.1, CLIP uses words in the WordNet synsets to search for images to construct the datasets used for pre-training, and 88% of all words in our datasets exist in WordNet, meaning that CLIP has seen them and how the visual concepts are demonstrated in images, but it cannot adapt the concepts to a

different visual domain, sketches, without fine-tuning from sketches in both categories.

5.2 Word Similarity Experiment



Figure 5.1: An example pair of sketches with different eyes that was presented to annotators. We do not pre-define a list of attributes and ask the annotators to determine whether the sketches are different in these attributes; instead, we highlight the parts in colors and implicitly prompt them to pick up on the differences. As expected, most annotations naturally include the size and shape differences. From the descriptions, we can extract pairs of *antonyms* used to indicate opposite visual concepts. In this case, the annotation was *large circular eyes* for the sketch on the left and *tiny solid eyes* for the one on the right.

In order to learn more about how word embeddings changed through fine-tuning, we use cosine similarity as a metric to measure how close two words are. During annotation, we present two sketches with contrasting parts (for example, a face with big eyes and another with small eyes, like Figure 5.1) to the annotators and ask them for descriptions of the parts. We do not explicitly tell them that the eyes are different in sizes, yet it is highly likely that they would provide descriptors that can indicate this visual difference. We extract all tuples of words from a pair of contrasting descriptions. In the case of Figure 5.1, one annotator provided the descriptions *large circular eyes* for the sketch on the left and *tiny solid eyes* for the other; from these two descriptions, we can extract 4 pairs of contrasting descriptors: (*large, tiny*), (*large, solid*), (*circular, tiny*), and (*circular, solid*). There are a total of 9,823 pairs, and the top 20 most frequent pairs are shown in Table 5.2. For example, the pair *small* and *large* are used 228 times by annotators to contrast certain parts in the two sketches across face and angel sketches. We consider two words in each pair as “antonyms”, because they are used to represent two *opposite* visual concepts in the two sketches. It is in cases like this that our common understanding of synonym and antonym falls short: (*large, tiny*) might align with most people’s idea of two words opposite in meaning, and (*large, solid*) do not appear as antithesis of each other, but the pair is used for the purpose of differentiating the sketches. Therefore, from a pragmatic perspective, these words are antonyms in our dataset.

word 1	word 2	#occurrences
small	large	228
circular	oval	142
wide	narrow	124
small	round	112
wide	small	108
oval	small	100
circular	small	94
oval	round	92
close	open	83
big	small	71
short	long	68
curve	wide	68
thin	wide	61
open	small	60
oval	wide	60
large	round	55
round	uneven	53
oval	open	51
open	circular	47
curve	small	46

Table 5.2: 20 pairs of descriptors most frequently used by annotators to differentiate parts in face and angel sketches. There are a total of 9823 pairs, and 7194, or 73.2%, pairs are used only once in the dataset.

	avg cosine sim
zero-shot	0.833
finetuned on face	0.752
finetuned on face + angel	0.691

Table 5.3: Average cosine similarity.

After fine-tuning CLIP to learn to contrast sketches with these descriptions, we expect that the similarity between contrasting words to decrease. Indeed, as shown in Table 5.3, we see a trend that as CLIP learns from more pairs of (sketch, part description) in our dataset, the word embeddings of contrasting words become dissimilar.

Across the entire dataset, there are $\binom{1450}{2}$, about 1 million, pairs of words, whether they are used in contrasting descriptions or not, and 99.9% of the pairs decreased in cosine similarity, from pre-trained CLIP to CLIP fine-tuned on face and angel sketches. One possible explanation is that we collected our data by presenting two sketches with contrasting features, but it could also be a general trend of CLIP fine-tuning on additional datasets, and we need future studies to give an definitive answer.

We only see increase in cosine similarity in 824 pairs of words. When we look at the top 20 pairs of words with the most increase in cosine similarity, we see that 18 pairs include the word *slash*, which appears only once in our entire dataset. The increase is less than 0.06, compared to a decrease of 0.65 for the pair that has the largest decrease in similarity. If we only look at pairs of words that are among the 250 most frequent words, pairs whose cosine similarity has increased the most are *arrow, smiley, circle, arrow, trilateral, triangular*. However, the increase is still too marginal, ≈ 0.02 , for us to conclude anything meaningful about the text embeddings of the fine-tuned CLIP model. Therefore, we cannot provide reasons for the marginal increase, and it is likely to be a chance event due to rare occurrences.

Chapter 6

Future Work



Figure 6.1: .

There are many creative use of language in our dataset, and in follow-up work, we want to investigate what categories of creative usage are there and quantify how many cases are in each category. Most of our current characterization of the dataset are based on looking through a few examples in the datasets (there are a total of 11K!). Since we did not put limits on what language annotators could use during data collection, such as specifying a fixed list of attributes for the annotators to choose from, the dataset contains a variety of ways of describing the same concept. For example, for a nose that looks similar to the one in Figure 6.1, annotators describe it as *hook-shaped*, *c-shaped*, *inverted j*, *curved*, etc. We believe that this case is not the only one, and we want to quantify the diversity. In this way, we can evaluate how a multi-modal embedding space, such as the CLIP joint embedding, captures these relationships. Is the feature vector of the sketch collinear with the word features of every one of these words?

So far, we have observed that face and angel sketches share many visual concepts. These could be related to length (*long, short*), size (*big, small*), geometry (*circular, triangular*), direction (*horizontal, vertical*), and many more. We want to evaluate more thoroughly which concepts are shared and which ones are unique to each category. In this way, we can understand better the challenges around generalizing CLIP to unseen categories, indicated by results in Section 5.1.

Appendix A

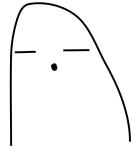
Designing the Requirement Section for the Prompt-Guided Sketch Text Dataset

An excerpt from an old version of the instructions, in which we tried to explain a single *step* of the annotation:

In each task, we show **1 prompt** from which we would like to get

1. A drawing containing **1 entity** that you think illustrates the prompt.
2. Text annotation for every “**component**” that makes up the entity. The word “component” is intentionally vague, and it depends on how you compose your drawing. For example, in the above example, the prompt is “smiley face”, and during the process of creating a “smiley face” entity, we used 4 components: a face, a left eye, a right eye, and a mouth. For each component, you can annotate with “face”, “left eye”, “right eye”, “mouth”, respectively; you can also annotate with more details describing the shapes of each component: “face that looks like an arc opening downwards”, “a left eye that is an arc”, “a right eye that looks exactly like the left eye”, “an arc-like mouth”. Try to use creative and descriptive languages. You can draw a component using multiple strokes.

We also need to come up with examples explaining each requirement. We select a few major sub-versions of the requirements resulted from circulating the interface within our lab.



(a) An example included in the first version of the requirements, explained in more details in item 3.

1. Do not draw entity that does not respond to the prompt. For example, given the prompt *Smiley Face*, the drawing should not contain irrelevant objects like a house.
2. Do not draw more than one entity that responds to the prompt. For example, One should not draw two *Smiley Face* entities, although each *Smiley Face* entity is good. However, you can draw multiple tree objects to illustrate the prompt *Forest*.
3. Do not draw entity that is ambiguous in terms of illustrating the prompt. For example, the drawing (Figure A.1a) looks more like a *Sad Face* than *Smiley Face*.
4. Do not draw one component that contains more information/content than what the annotation for that component describes. (A counterexample is illustrated in Figure A.1b.)
5. Do not split the drawing of a component into multiple steps, unless you can annotate each step separately. (A counterexample is illustrated in Figure A.1c.)
6. Do not annotate a component more than once.

1. Draw *one* item at a time and provide its corresponding annotation. For example, the text annotation says “left eye”, but two items are drawn: a left eye and a right eye.
2. The annotation should describe its corresponding item in the drawing *entirely*.
3. The annotation should name the item.
4. Desired properties of good drawings:

- Contain as many items as possible, but be sure that they all contribute to illustrating the prompt. For example, draw more than just two eyes for a face.
- Use shapes creatively. For example, draw a triangle for the left eye, and annotate accordingly with “triangular left eye that shows suspicion”.

Bad		
Step 1		
Prompt	Canvas	Component Annotation
Smiley Face		Face that looks like a curve opening downward
Step 2		
Prompt	Canvas	Component Annotation
Smiley Face		Face that looks like a curve opening downward Left eye
The drawing contains information about both left and right eye, which is more than what the annotation describes.		

(b) Unaligned drawing and text description.

Bad		
Step 1		
Prompt	Canvas	Component Annotation
Eye		Eye socket
Step 2		
Prompt	Canvas	Component Annotation
Eye		Eye socket Eye socket
Two "Eye socket" annotations would indicate that the drawing contains two components that are "Eye socket", but it is one "Eye-socket" split into 2 steps.		

(c) An example of misalignment: text description *overflow* into multiple steps.

Figure A.1: Screenshots of counterexamples used in first version of the requirements in Version 1.

5. Desired properties of good annotations:

- Use descriptive languages. For example, “a left eye that looks an arc facing downward”.
- Include the intention/purpose of drawing an item. Explain in the annotation reasons for drawing the item. For example, “thumbs-up next to the face that really shows how happy

the face is”.

1. Each drawing should contain at least 2 steps.
2. Annotation of each step should include at least the name of the drawn object(s).
3. If draw multiple copies of the same object, draw each object in a separate step and give different annotations by using, for example, cardinal or ordinal numbers. (An example shown in Figure A.2)
4. Differentiate between plural and singular forms.
5. The name of the whole should not be used for its parts. (An example shown in Figure A.2)
6. The word ”right” always refers to this side: \Rightarrow

Prompt	Canvas	Annotation										
A house		<div style="display: flex; align-items: center;"> <input type="text" value="Annotate an item..."/> <button>Add</button> </div> <table border="1" style="width: 100%; border-collapse: collapse;"> <tr> <td style="padding: 2px;">triangle</td> <td style="padding: 2px;"><button>Delete</button></td> </tr> <tr> <td style="padding: 2px;">rectangle</td> <td style="padding: 2px;"><button>Delete</button></td> </tr> <tr> <td style="padding: 2px;">rectangle with cross inside</td> <td style="padding: 2px;"><button>Delete</button></td> </tr> <tr> <td style="padding: 2px;">rectangle with cross inside</td> <td style="padding: 2px;"><button>Delete</button></td> </tr> <tr> <td style="padding: 2px;">rectangle</td> <td style="padding: 2px;"><button>Delete</button></td> </tr> </table>	triangle	<button>Delete</button>	rectangle	<button>Delete</button>	rectangle with cross inside	<button>Delete</button>	rectangle with cross inside	<button>Delete</button>	rectangle	<button>Delete</button>
triangle	<button>Delete</button>											
rectangle	<button>Delete</button>											
rectangle with cross inside	<button>Delete</button>											
rectangle with cross inside	<button>Delete</button>											
rectangle	<button>Delete</button>											
Problematic Step	Reason											
1,2,3,4,5	Violates: <ul style="list-style-type: none"> • Annotation of each step should include at least the name of the drawn object(s). 											
3,4	Violates: <ul style="list-style-type: none"> • Give different annotations if draw multiple copies of the same object. For example, by adding "left window that is a cross inside a rectangle" 											

Figure A.2: Screenshots of counterexamples used in third version of the requirements.

Bibliography

- Allado-McDowell, K., & Okojie, I. (2020). *Pharmako-ai*. Ignota.
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... Amodei, D. (2020). *Language models are few-shot learners*. arXiv. Retrieved from <https://arxiv.org/abs/2005.14165> doi: 10.48550/ARXIV.2005.14165
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... Houlsby, N. (2020). *An image is worth 16x16 words: Transformers for image recognition at scale*. arXiv. Retrieved from <https://arxiv.org/abs/2010.11929> doi: 10.48550/ARXIV.2010.11929
- Eitz, M., Hays, J., & Alexa, M. (2012). How do humans sketch objects? *ACM Trans. Graph. (Proc. SIGGRAPH)*, 31(4), 44:1–44:10.
- Frans, K., Soros, L. B., & Witkowski, O. (2021). *Clipdraw: Exploring text-to-drawing synthesis through language-image encoders*. arXiv. Retrieved from <https://arxiv.org/abs/2106.14843> doi: 10.48550/ARXIV.2106.14843
- Gal, R., Patashnik, O., Maron, H., Chechik, G., & Cohen-Or, D. (2021). *Stylegan-nada: Clip-guided domain adaptation of image generators*. arXiv. Retrieved from <https://arxiv.org/abs/2108.00946> doi: 10.48550/ARXIV.2108.00946
- Ge, S., Goswami, V., Zitnick, C. L., & Parikh, D. (2020). *Creative sketch generation*. arXiv. Retrieved from <https://arxiv.org/abs/2011.10039> doi: 10.48550/ARXIV.2011.10039
- Ha, D., & Eck, D. (2017). *A neural representation of sketch drawings*.
- Li, K., Pang, K., Song, J., Song, Y.-Z., Xiang, T., Hospedales, T. M., & Zhang, H. (2018). *Universal perceptual grouping*. arXiv. Retrieved from <https://arxiv.org/abs/1808.02312> doi: 10.48550/ARXIV.1808.02312
- Nichol, A., Dhariwal, P., Ramesh, A., Shyam, P., Mishkin, P., McGrew, B., ... Chen, M. (2021). *Glide: Towards photorealistic image generation and editing with text-guided diffusion models*. arXiv. Retrieved from <https://arxiv.org/abs/2112.10741> doi: 10.48550/ARXIV.2112.10741

- Patashnik, O., Wu, Z., Shechtman, E., Cohen-Or, D., & Lischinski, D. (2021). *Styleclip: Text-driven manipulation of stylegan imagery*. arXiv. Retrieved from <https://arxiv.org/abs/2103.17249> doi: 10.48550/ARXIV.2103.17249
- Qi, Y., & Tan, Z.-H. (2019). Sketchsegnet+: An end-to-end learning of rnn for multi-class sketch semantic segmentation. *IEEE Access*, 7, 102717-102726. doi: 10.1109/ACCESS.2019.2929804
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... Sutskever, I. (2021). *Learning transferable visual models from natural language supervision*. arXiv. Retrieved from <https://arxiv.org/abs/2103.00020> doi: 10.48550/ARXIV.2103.00020
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language models are unsupervised multitask learners..
- Ramesh, A., Dhariwal, P., Nichol, A., Chu, C., & Chen, M. (2022). *Hierarchical text-conditional image generation with clip latents*. arXiv. Retrieved from <https://arxiv.org/abs/2204.06125> doi: 10.48550/ARXIV.2204.06125
- Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., ... Sutskever, I. (2021). *Zero-shot text-to-image generation*. arXiv. Retrieved from <https://arxiv.org/abs/2102.12092> doi: 10.48550/ARXIV.2102.12092
- Sangkloy, P., Burnell, N., Ham, C., & Hays, J. (2016, jul). The sketchy database: Learning to retrieve badly drawn bunnies. *ACM Trans. Graph.*, 35(4). Retrieved from <https://doi.org/10.1145/2897824.2925954> doi: 10.1145/2897824.2925954
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... Polosukhin, I. (2017). *Attention is all you need*. arXiv. Retrieved from <https://arxiv.org/abs/1706.03762> doi: 10.48550/ARXIV.1706.03762
- Yu, Q., Liu, F., Song, Y.-Z., Xiang, T., Hospedales, T. M., & Loy, C. C. (2016). Sketch me that shoe. In *2016 ieee conference on computer vision and pattern recognition (cvpr)* (p. 799-807). doi: 10.1109/CVPR.2016.93